

Welcome! While we wait...



Workshop material


<https://github.com/weaviate-tutorials/scalable-ai-workshop>



Tell us a little bit about yourself!

Short survey:

<https://www.menti.com/al31ja5gtijw>





How to build scalable AI applications with a vector database



JP Hwang
Developer Educator

Agenda

- **Recap: Vector DBs & RAG**
- **Challenges of scalability**
 - **Work with a local cluster**
 - **Apply practical solutions**
- **Challenges of reliability**

Let's get started

While we wait, go to:

- **<https://github.com/weaviate-tutorials/scalable-ai-workshop>**

Start on the README instructions


- **Step 1**
- **Step 2**



Recap: Vector DBs & RAG


Large Language Model

Can “hallucinate”





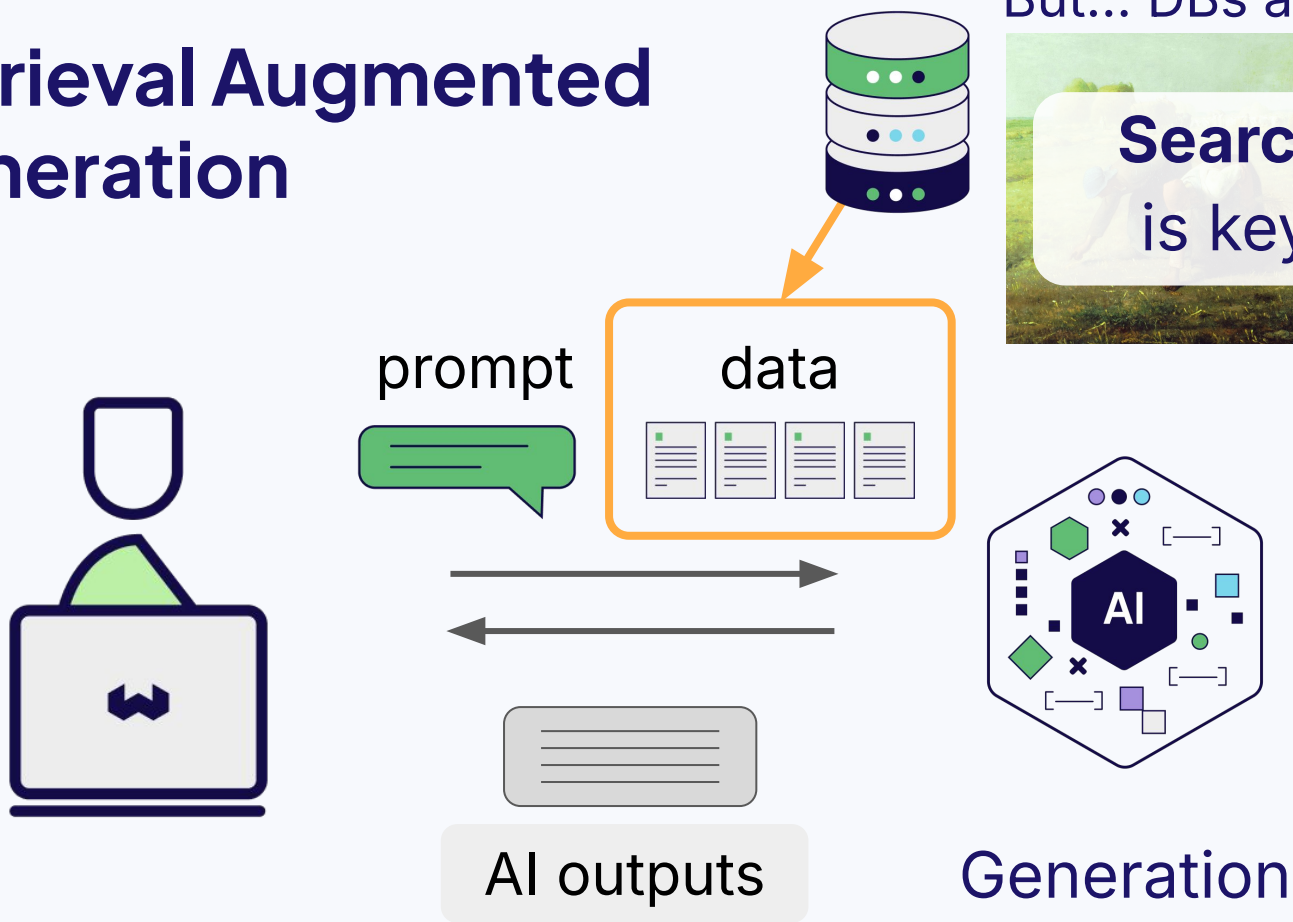
Retrieval Augmented Generation





Retrieval Augmented Generation

But... DBs are big!





Workshop / Live coding

Challenges of scale


Scale: Considerations

- **Object count:** Memory & storage
- **User count:** Data management & compliance
- **Server load:** Distribute load

Managing resource requirements


Scale: Solutions

Growth




Scale: Solutions

Growth




Scale: Solutions

Growth





Scale: Solutions

Growth



Scale: Solutions

Growth



- Single point of failure
- Efficiency
- Upgrades
- Costs



🔍 1024 gib



19 matches

Instance name ▾	On-Demand hourly rate ▲	vCPU ▾	Memory ▾	Storage ▾
x2gd.metal	\$5.344	64	1024 GiB	2 x 1900 SSD
x2gd.16xlarge	\$5.344	64	1024 GiB	2 x 1900 SSD
hpc6id.32xlarge	\$5.70	64	1024 GiB	4 x 3800 NVMe SSD
x2iedn.8xlarge	\$6.669	32	1024 GiB	1 x 950 NVMe SSD
x2idn.16xlarge	\$6.669	64	1024 GiB	1 x 1900 NVMe SSD

AWS Spot pricing, Nov 2024



1024 gib

19 matches

Instance name

On-Demand hourly rate

Storage

x2gd.metal

\$5.344

1900 SSD

x2gd.16xlarge

hpc6id.32xlarge

x2iedn.8xlarge

x2idn.16xlarge



$120 * 365 = \$43,800 / \text{year!}$

AWS Spot pricing, Nov 2024



Vector indexing options




Scale: Solutions

Improve efficiency - indexing

Scale: Solutions

Improve efficiency - indexing

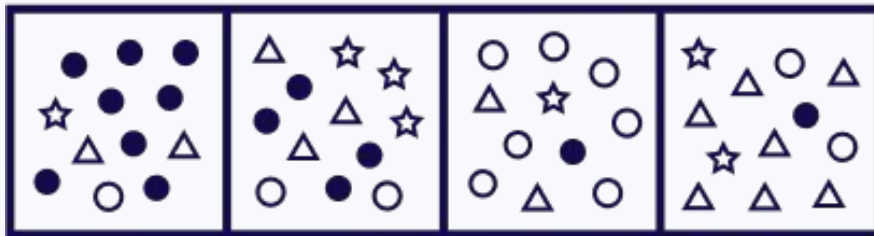
- **HNSW** index (default)



Scale: Solutions

Improve efficiency - indexing


- **HNSW** index (default)
- **Flat** index



Scale: Solutions

Indexes - comparison

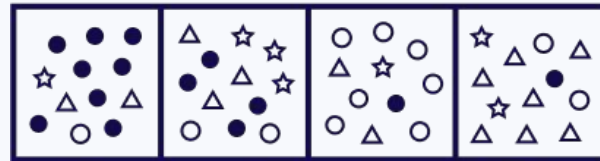
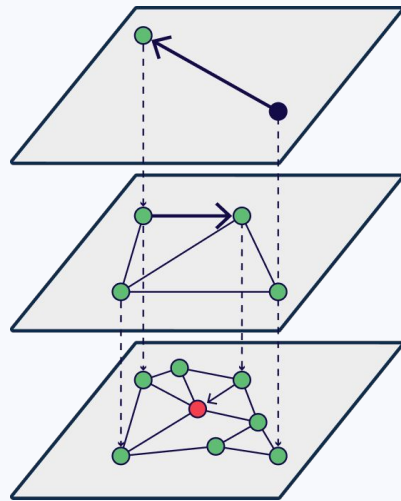
- **HNSW**: fast + scalable
- **Flat**: tiny footprint; ~100k objs



Scale: Solutions

How to choose?


- **Start with HNSW**
(Tune speed / size / accuracy)
- Multi-tenancy?
 - Try **dynamic**



Scale: Solutions

Improve efficiency - indexing

- **HNSW** index (default)
- **Flat** index
- **Dynamic** index
 - Flat → HNSW @ threshold

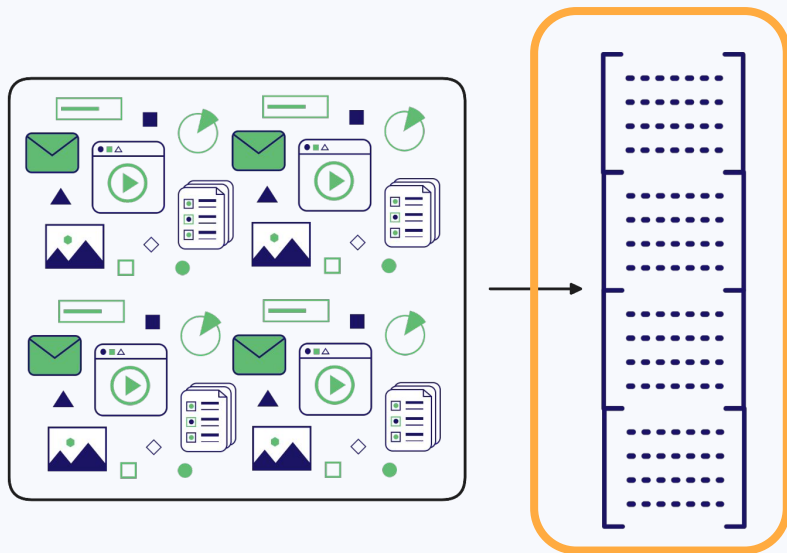




Vector quantization

Scale: Solutions

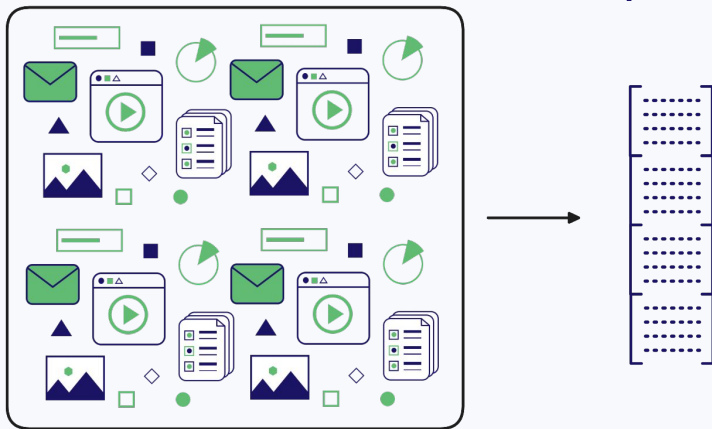
Improve efficiency



Scale: Solutions

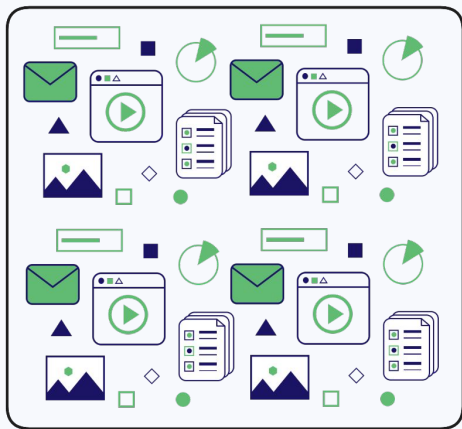
Improve efficiency

Compression




Scale: Solutions


Improve efficiency



Compression

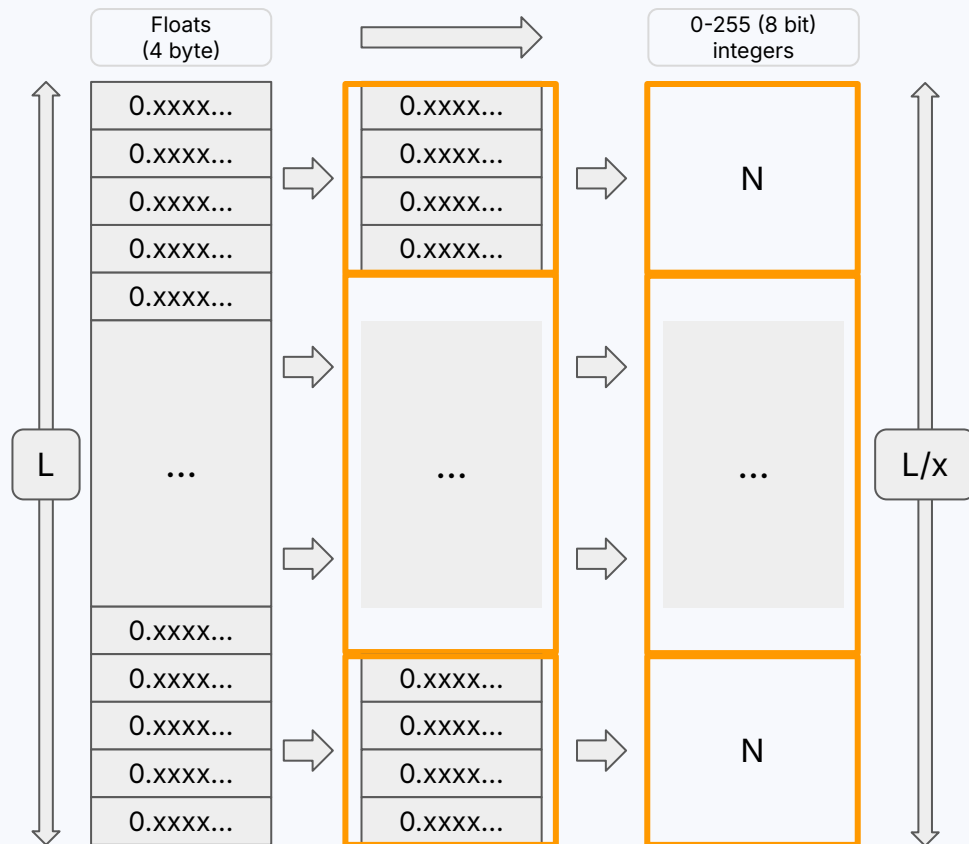


Product quantization



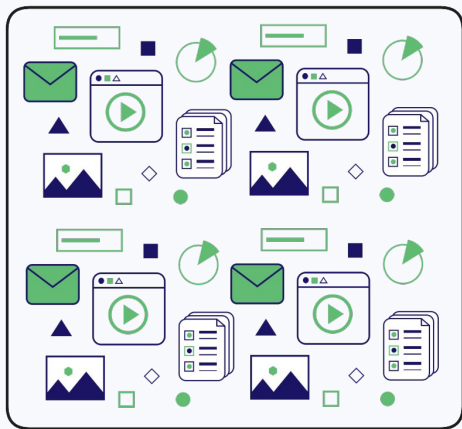
Customisable compression

(e.g. 128 floats → 32 bytes: 16x)

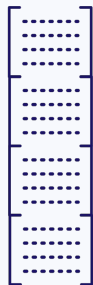


Scale: Solutions

Improve efficiency



Compression



Binary quantization

$[-0.1324..., -0.9253...,$
 $0.2389...,$
 \dots
 $0.3249..., 0.2390..., -0.4823...]$

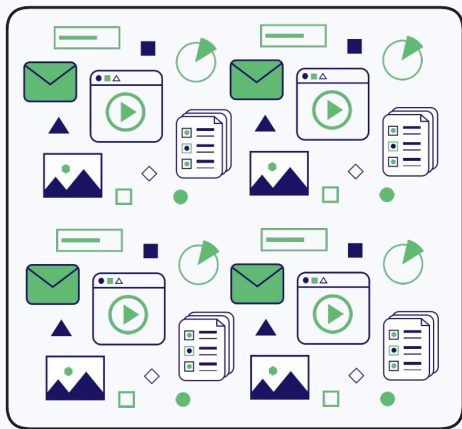


$[0, 0, 1, 0, 1, 1,$
 \dots
 $0, 1, 0, 1, 1, 0]$

$n \text{ floats} \rightarrow n \text{ bits}$
(32x reduction)

Scale: Solutions

Improve efficiency



Compression



Scalar quantization

$[-0.1324..., -0.9253...,$
 $0.2389...,$
 \dots
 $0.3249..., 0.2390..., -0.4823...]$




$[104, 12, 138,$
 \dots
 $152, 138, 47]$

$n \text{ floats} \rightarrow n \text{ ints}$
(4x reduction)



Example: Index type & quantization vs memory footprint



HNSW: ~6GB

**HNSW + BQ/PQ:
~2GB**

**Flat (+BQ):
~1GB**



Instance type

vCPU

32

Memory

\$58k / year

vs.

\$23k / year

View



< 1 2 >

Instance name

Demand
hourly rate

vCPU

Memory

Storage

Network
performance

x2iedn.8xlarge

\$6.669

32

1024 GiB

1 x 950 NVMe
SSD

25 Gigabit

x1e.8xlarge

\$6.672

32

976 GiB

1 x 960 SSD

Up to 10 Gigabit

x2gd.8xlarge

\$2.672

32

512 GiB

1 x 1900 SSD

12 Gigabit

AWS Spot pricing, Nov 2024


Scale: Solutions

BQ / PQ / SQ compression

Search **quality** mitigated by over-fetching & rescoring

When to use which?

- Generally, try PQ first
- BQ: model-specific






Multi-tenancy


Scale: Solutions

End user growth



Scale: Solutions

End user growth




Challenges faced:

- Performance
- Data isolation
- Compliance

Developed: **Multi-tenancy**

Scale: Solutions

End user growth




Multi-tenancy implementation

- 1000s per node
 - Isolated
 - Active/inactive/offloaded tenants
- tenants

Scale: Solutions

End user growth




Multi-tenancy implementation

- 1000s per node
- **Isolated** (easy deletion & compliance)
- Active/inactive/offloaded tenants

Scale: Solutions

End user growth



Multi-tenancy implementation

- 1000s per node
- Isolated
- **Active/inactive/offloaded tenants** (efficient)




Replication

Reliability: Considerations

- **Robust to errors:** Ensure consistency
- **Downtime:** Reduce disruption
- **Backups:** In case of emergency


Reliability: Solutions

Happy days




Reliability: Solutions

(Less) Happy days



Reliability: Solutions

(Less) Happy days




Downtime = inevitable

Reliability: Solutions

Provide redundancy


Replication



Reliability: Solutions

Provide redundancy


Replication



Reliability: Solutions

Provide redundancy

Replication



Reliability: Solutions


Provide redundancy

Node-level downtime:

- inevitable


System-level downtime:

- avoidable!




Reliability: Solutions

Sharding + replication




Reliability: Solutions

Sharding + replication




Upgrade: Solutions

Sharding + replication



Upgrade: Solutions

Sharding + replication




Upgrade: Solutions

Provide redundancy

"Just works"


for the end user -
with the latest
versions



Scaling out

Scale: Solutions

Deal with growth






Scale: Solutions

Deal with growth

To scale:
Add more nodes




Billion scale
since 2022!



Scale: Solutions

Deal with growth

To scale:
Add more nodes




Billion scale
since 2022!



Scale: Solutions

Deal with growth

To scale:
Add more nodes



Billion scale
since 2022!



Available Solutions

Scaling up

Scaling out

Index options

Quantization

Multi-tenancy
(+ tenant states)

Replication



Thank you

Connect with us!



weaviate.io



[weaviate/weaviate](https://github.com/weaviate/weaviate)



[weaviate_io](https://twitter.com/weaviate_io)




BONUS SLIDES!





Developer **workflow**


Ingest data




Ingest data



Ingest data




Ingest data

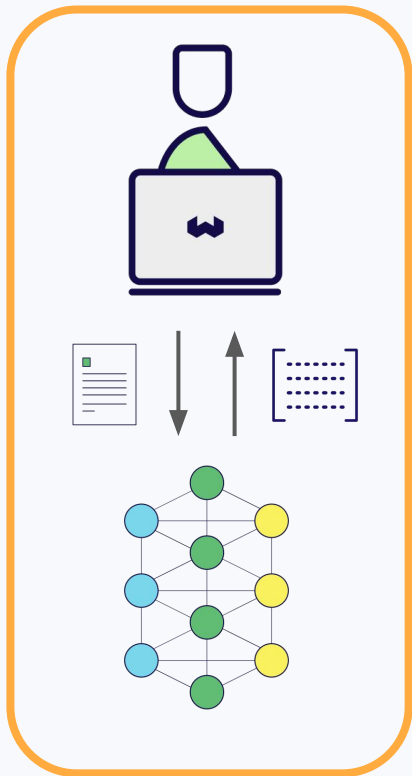


Where do they come from??

Obtain Embeddings



Obtain Embeddings



```
import cohere


co = cohere.Client("<<cohere_api_key>>")

response = co.embed(
    texts=["<YOUR_TEXT_INPUTS>"],
    model="embed-english-v3.0",
    input_type="search_document"
)

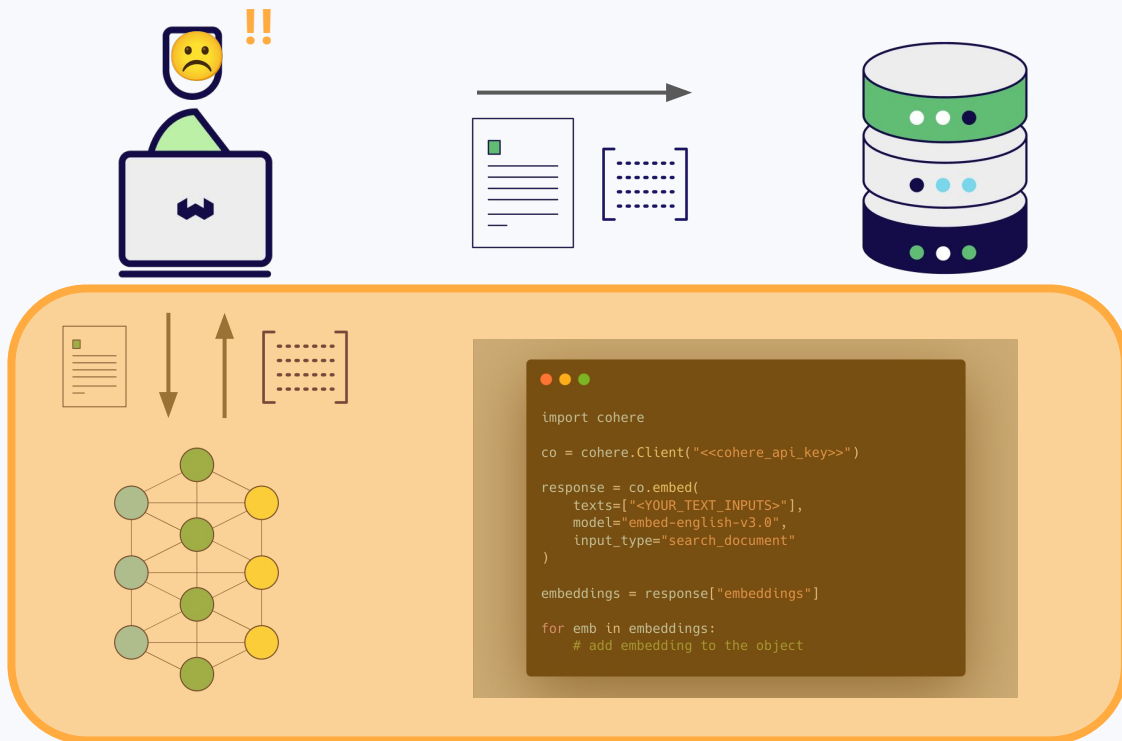
embeddings = response["embeddings"]

for emb in embeddings:
    # add embedding to the object
```

Obtain Embeddings




Overhead




Ingest data



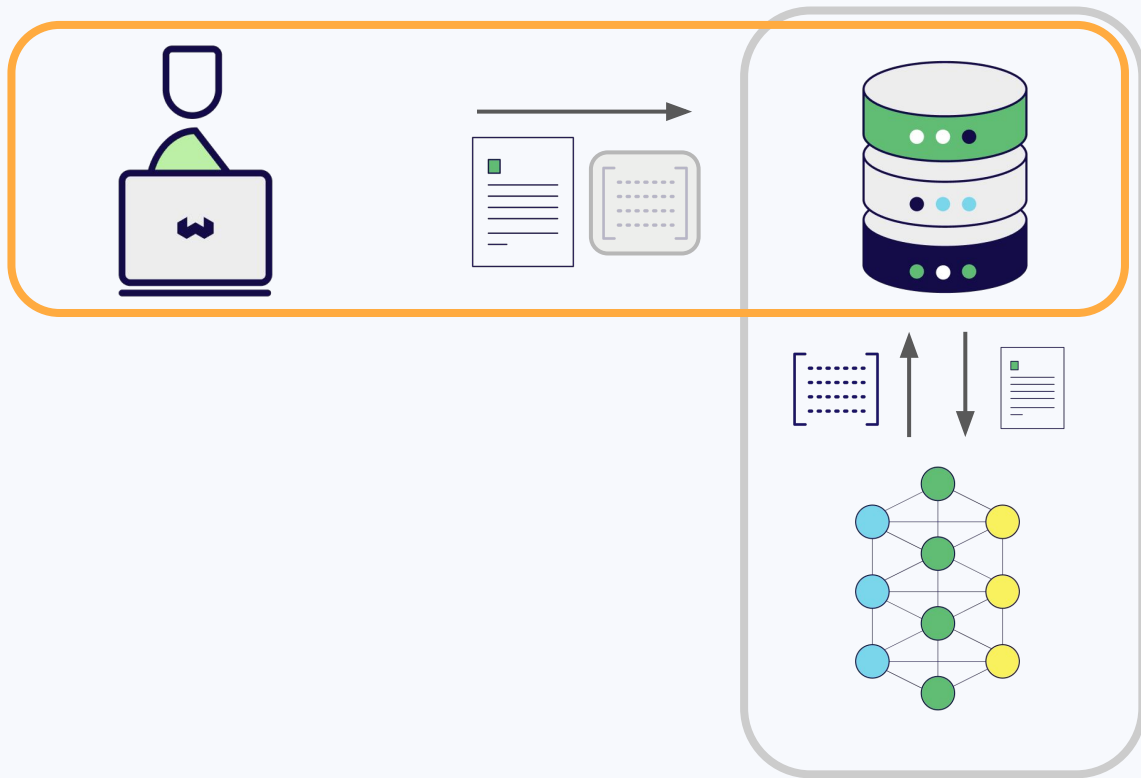
Have Weaviate obtain embeddings



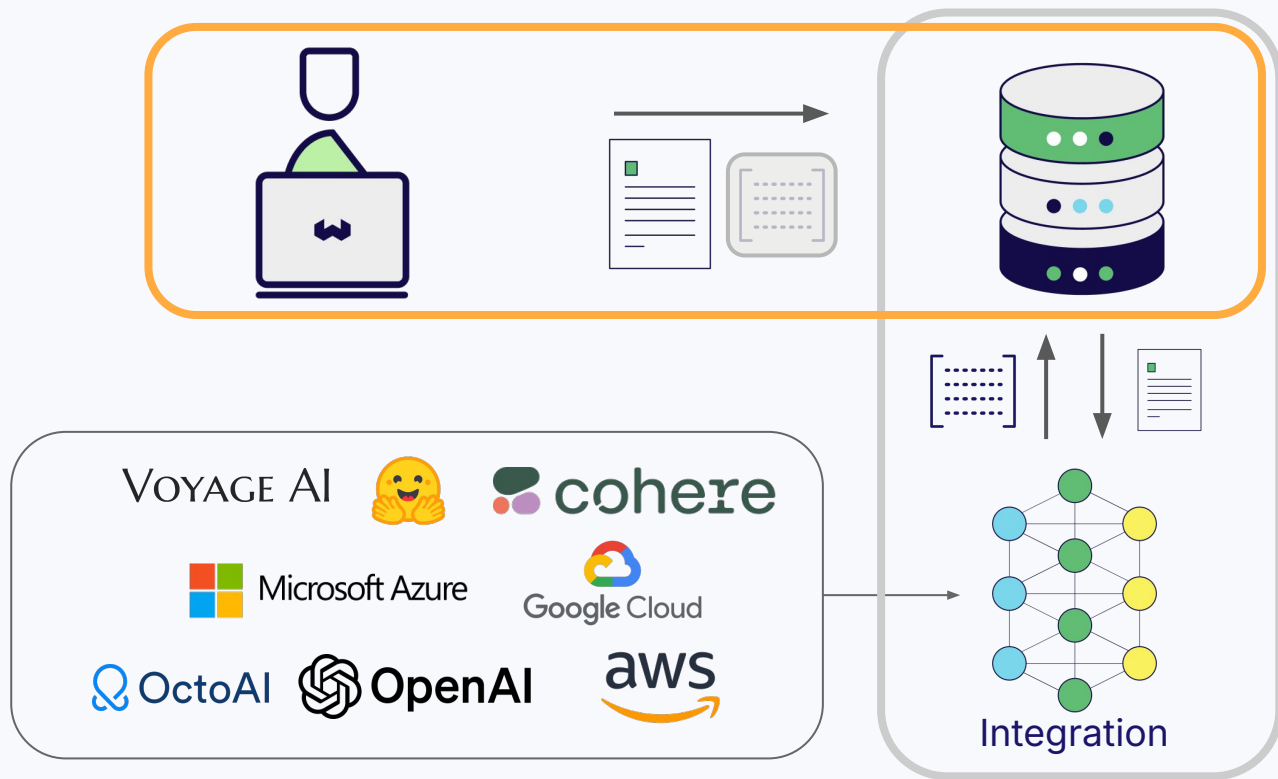
Ingest data (embeddings in background)




Ingest data (embeddings optional)




Ingest data (with provider integrations)



Simplified Experience




User
experience



Under the
hood



Simplified Experience




User
experience


```
collection = client.collections.get("YourCollection")

with collection.batch.dynamic() as batch:
    for data_row in data_rows:
        batch.add_object(
            properties=data_row,
        )
```



Query




Query



Query

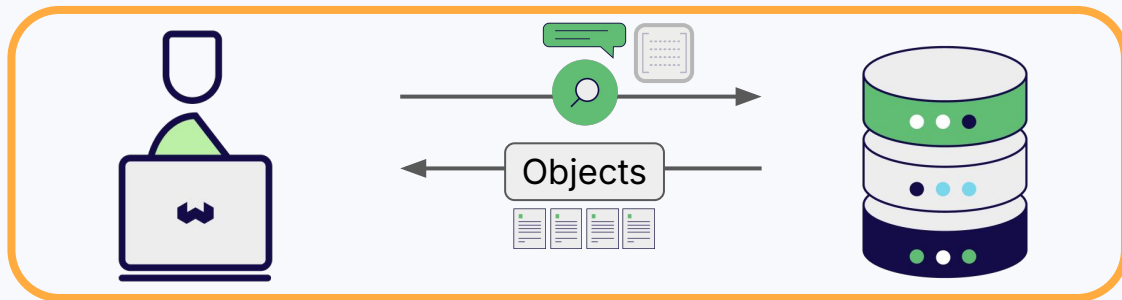


Query

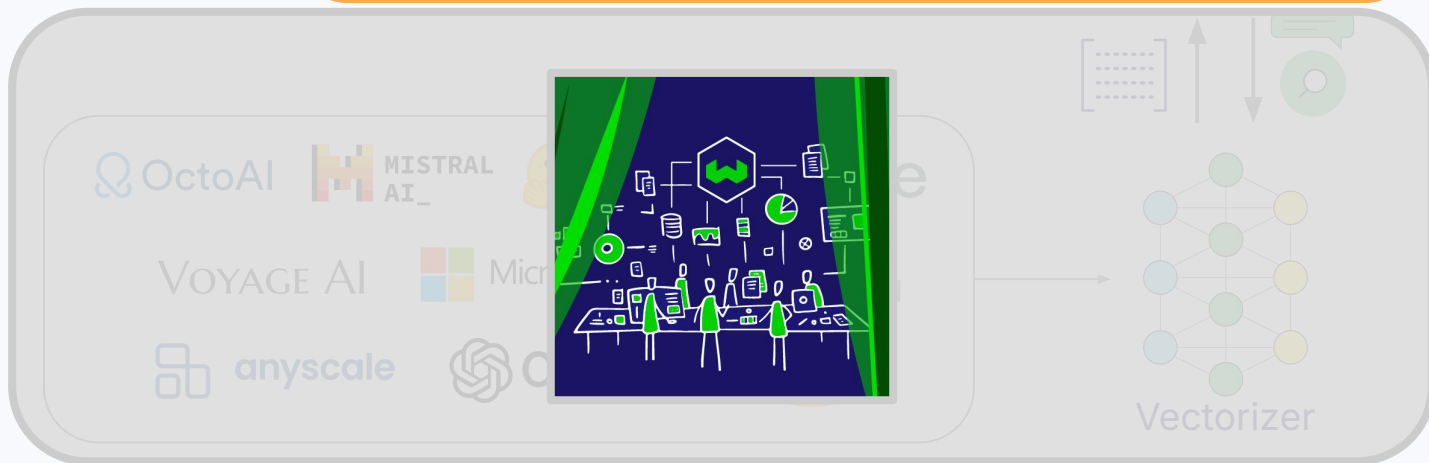




Query




User
experience




Under the
hood


RAG



RAG




RAG



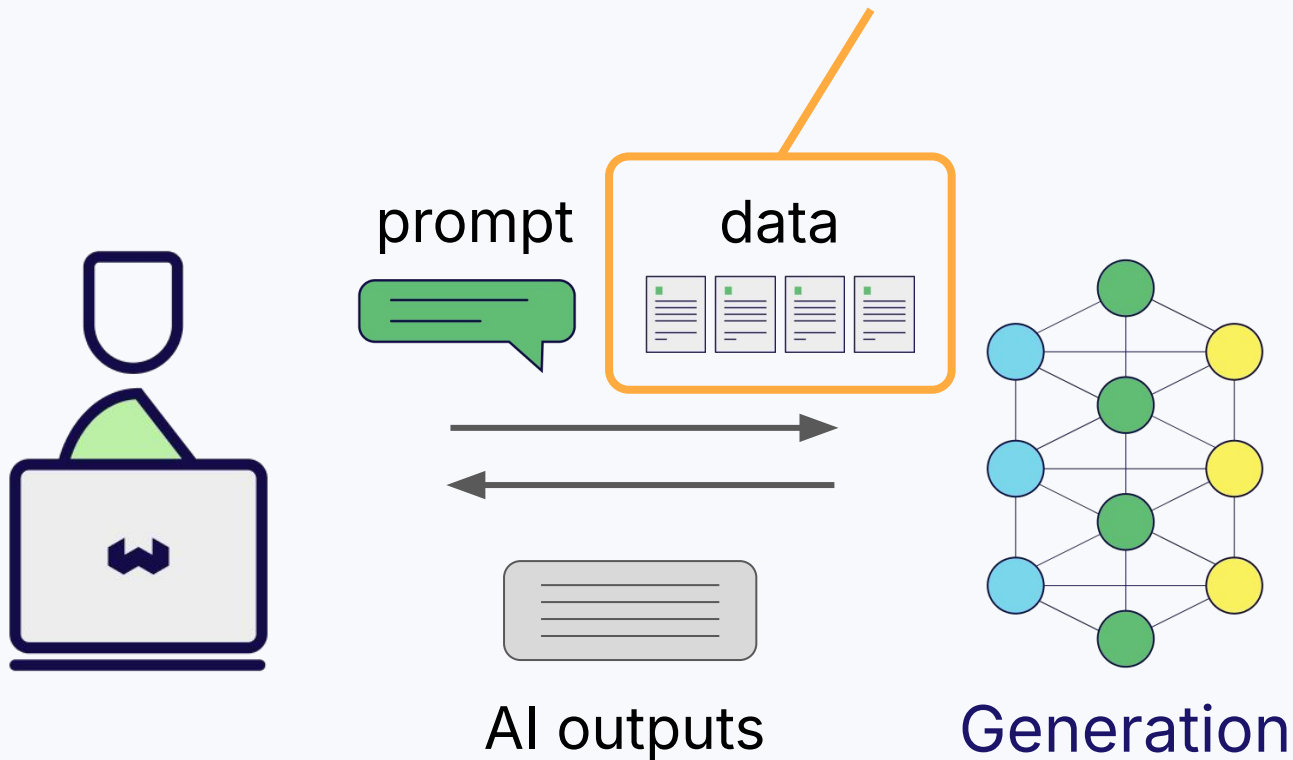
RAG

Vector search is key




RAG

Vector search: relevance search




RAG

Why **vector search** is hot




RAG






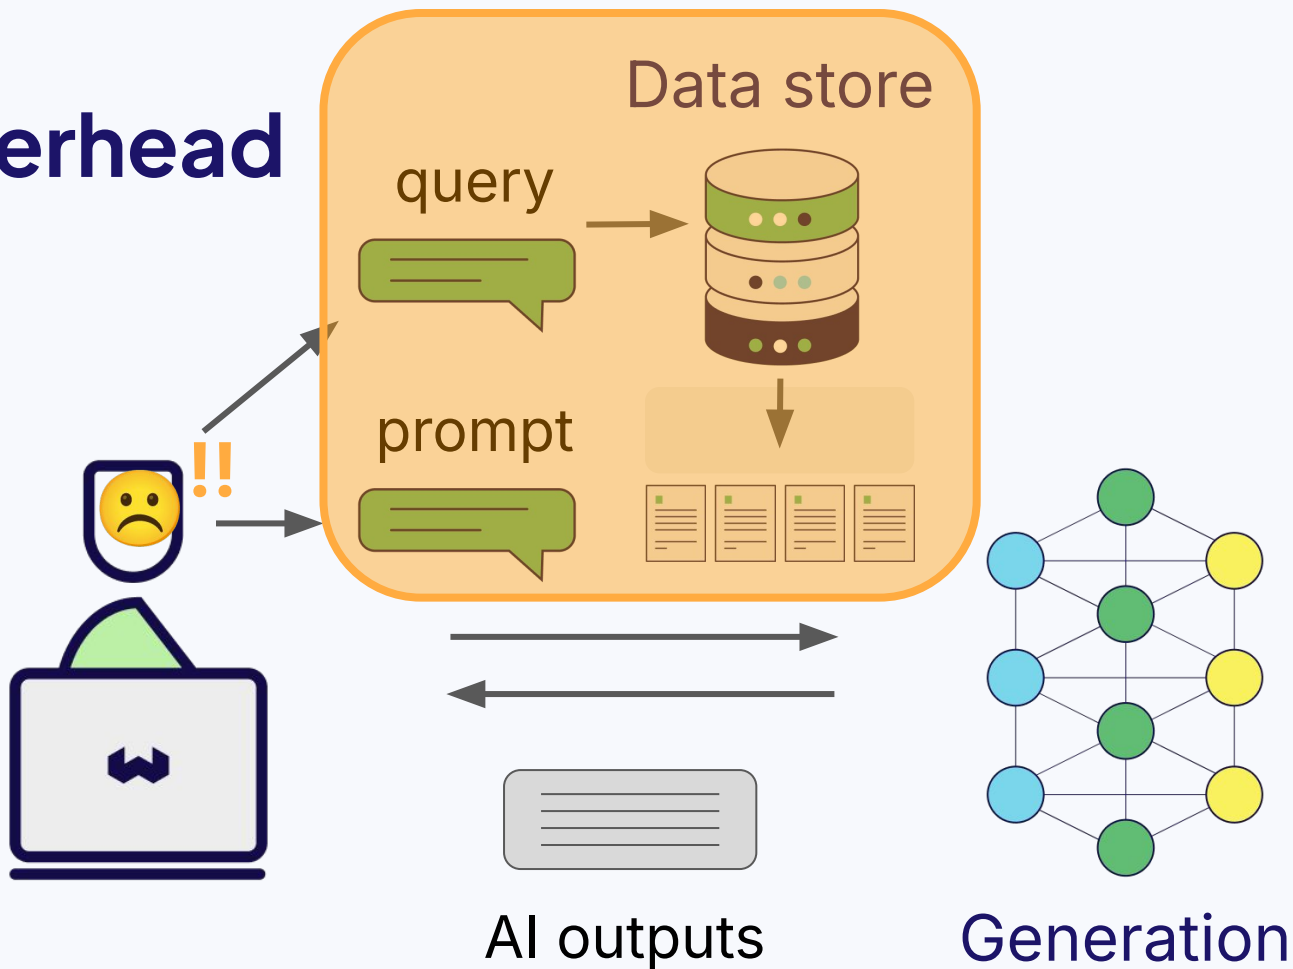
RAG




RAG




Overhead




RAG



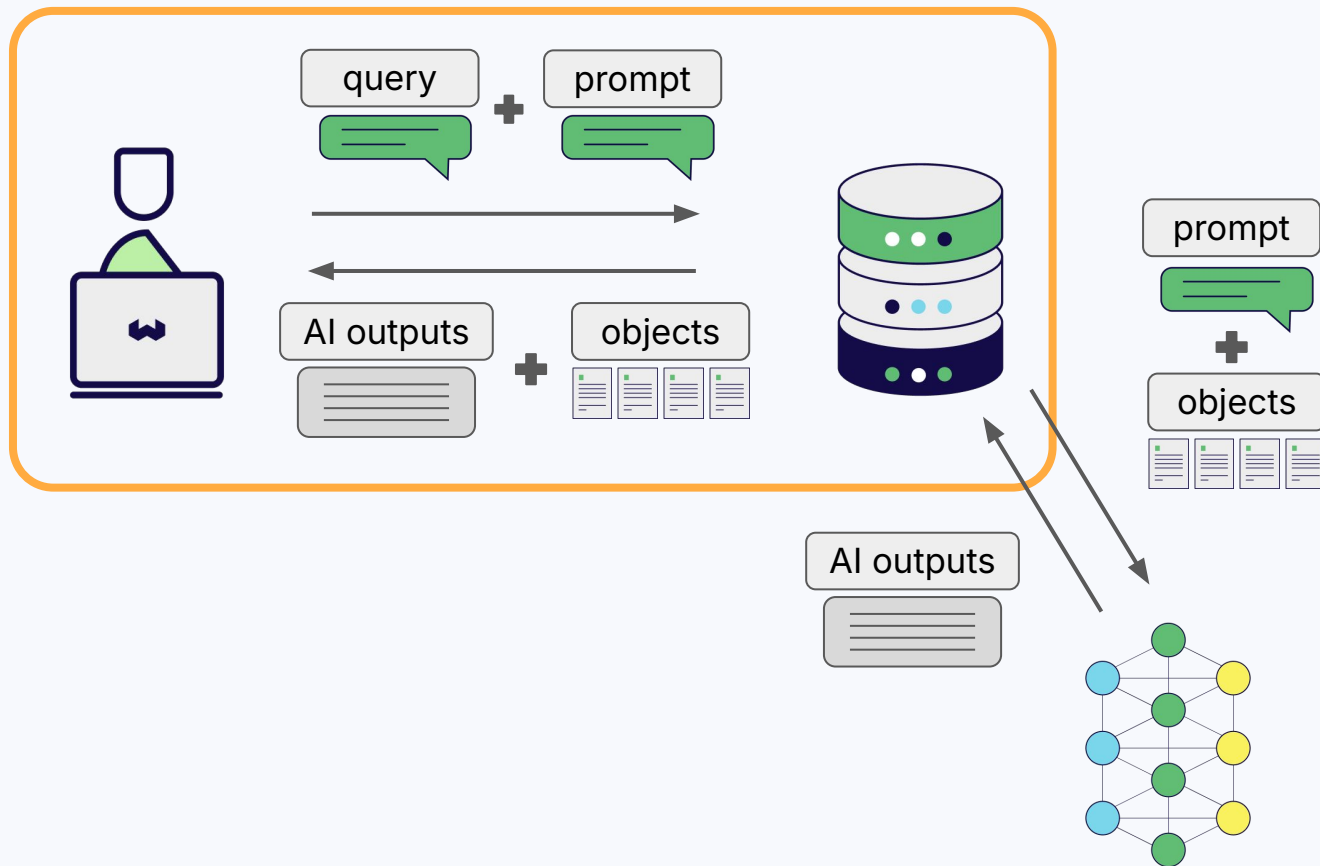
RAG



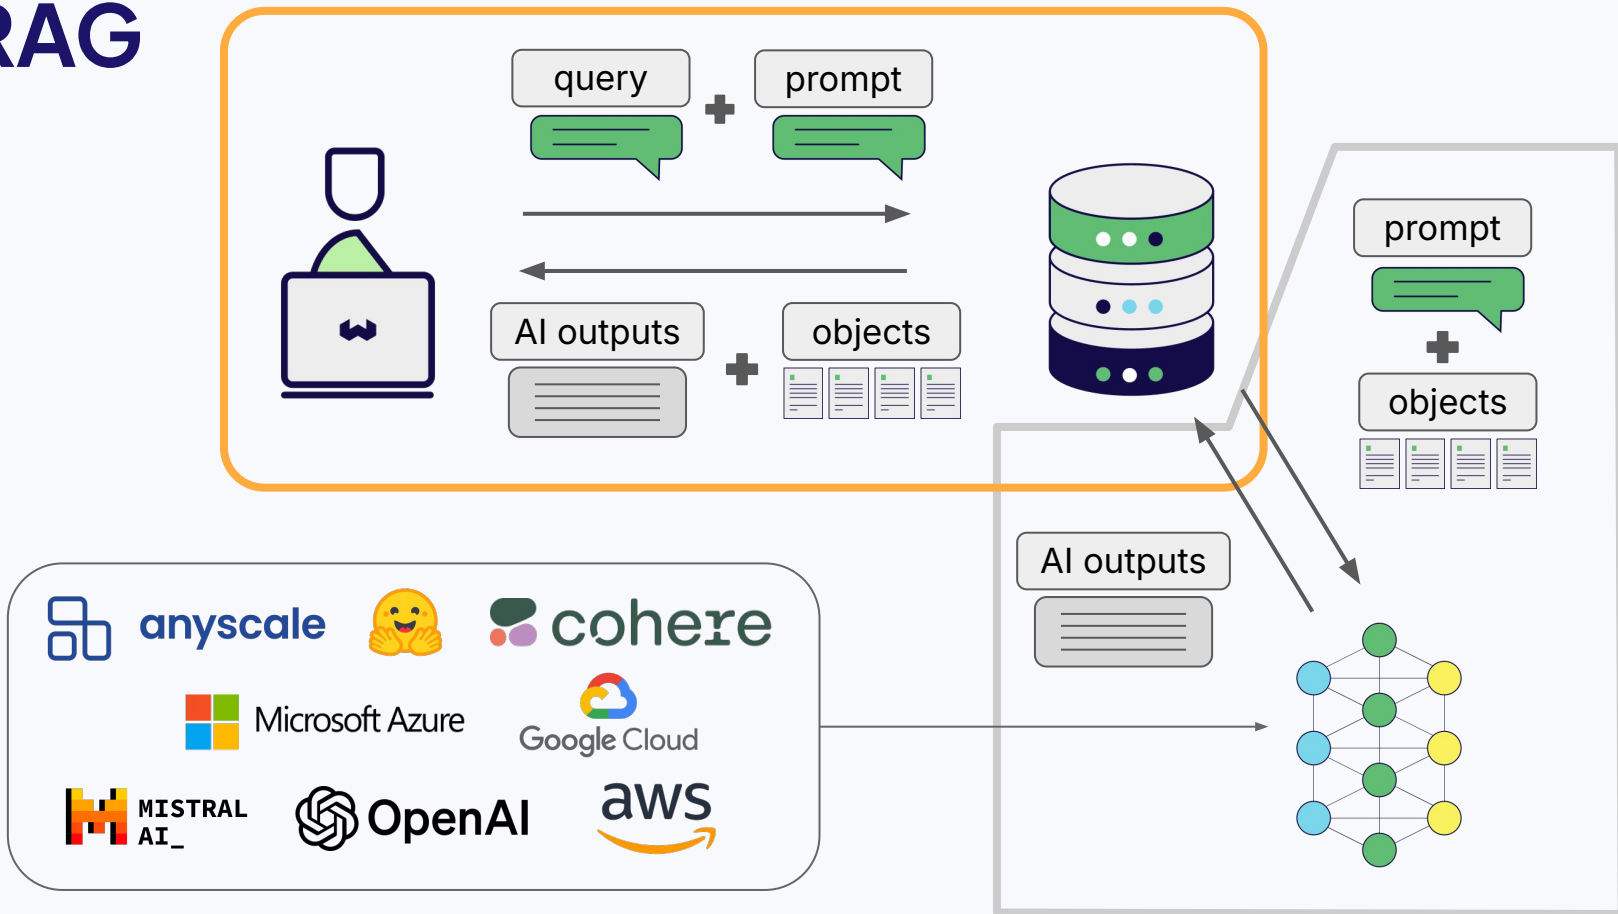
RAG



RAG



RAG



Simplified Experience

```
collection = client.collections.get("YourCollection")
response = collection.generate.hybrid(
    query="Multi-tenancy in Weaviate"
    limit=3,
    grouped_task="Explain how multi-tenancy works"
)

print(response.generated)
```

User
experience
(RAG)

RAG

User
experience

